

In [1]:

```
import warnings
warnings.filterwarnings("ignore")
from sklearn.datasets import load_boston
from random import seed
from random import randrange
from csv import reader
from math import sqrt
from sklearn import preprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from prettytable import PrettyTable
from sklearn.linear_model import SGDRegressor
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
```

In [77]:

```
X = load_boston().data
Y = load_boston().target
```

In [78]:

```
scaler = preprocessing.StandardScaler().fit(X)
X = scaler.transform(X)
```

In [79]:

```
clf = SGDRegressor()
clf.fit(X, Y)
print(mean_squared_error(Y, clf.predict(X)))
```

22.002611096122383

In [80]:

```
Y[0:10]
```

Out[80]:

```
array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9])
```

In [81]:

```
clf.predict(X)[0:10]
```

Out[81]:

```
array([30.11541469, 24.80979717, 30.33241595, 28.38883431, 27.73980678,
       25.04009414, 22.83449853, 19.40492534, 11.48135014, 18.7864729 ])
```

implementation of SGD manually:

In []:

In [82]:

```
## loading the data
data= pd.DataFrame(load_boston().data)
```

```
data.shape
```

```
Out[82]:
```

```
(506, 13)
```

```
In [ ]:
```

```
In [83]:
```

```
##initialize the X and Y
X=data
Y=load_boston().target
X.head()
```

```
Out[83]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [84]:
```

```
## split the data in train and test data set
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.33, random_state = 5)
### print shape of the train and test dataset
print('x_train data shape: ',x_train.shape)
print('y_train data shape: ',y_train.shape)

print('x_test data shape: ',x_test.shape)
print('y_test data shape: ',y_test.shape)
```

```
x_train data shape: (339, 13)
y_train data shape: (339,)
x_test data shape: (167, 13)
y_test data shape: (167,)
```

```
In [85]:
```

```
#preprocessing
## apply StandardScaler to scale the x_train and x_test data
scaler = preprocessing.StandardScaler().fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)
## make dataframe of standardized train and test data
x_train = pd.DataFrame(x_train)
x_test = pd.DataFrame(x_test)

# add last column as target column in dataframe
x_train['cost'] = np.array(y_train)
x_test['cost'] = np.array(y_test)
```

```
In [86]:
```

```
x_train.head()
```

```
Out[86]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.911839	0.502419	1.072305	0.256978	1.633548	0.486034	0.962774	0.823477	1.655334	1.552100	0.808078	0.842950	1.523203

	0	1	2	3	4	5	6	7	8	9	10	11	12
1	0.411727	0.502419	1.129795	0.256978	0.552451	1.028078	0.668619	0.183274	0.871371	0.802704	0.304174	0.427436	0.995240
2	0.124583	0.502419	1.072305	0.256978	1.441946	3.913414	0.725324	1.075955	1.655334	1.552100	0.808078	0.053353	0.765646
3	0.406208	0.839388	0.901940	0.256978	1.083710	0.097426	0.515087	1.600509	0.411970	0.624310	0.860301	0.152292	0.184576
4	0.021742	0.502419	1.072305	0.256978	1.398401	0.123238	0.743044	0.605107	1.655334	1.552100	0.808078	0.365116	0.301538

In []:

In []:

In [94]:

```
### initialization of the parameters
w = np.zeros(13) # initial weight
n = 10000 # no of iterations
r = 0.01 # learning_rate
k = 40 # batch_size
b = 0 #consider zero_intercept intitially
```

In []:

In [95]:

```
for i in range(1,n):
    batch = x_train.sample(k) ## choose random k samples from x_train
    x_tmp = np.array(batch.drop('cost',axis=1))
    y_tmp = np.array(batch['cost'])
    w_temp = np.zeros(13) ## intialize w_temp with 0
    b_temp = 0
    for j in range(1,k):
        val1 = -2 * x_tmp[j]
        val2 = y_tmp[j]
        #print(w.shape,x[z].shape)
        val3 = np.dot(w,x_tmp[j])
        val5 = (val2 - val3)
        val4 = (val1 * val5)
        w_temp= w_temp + val4
        b_temp = (-2)*( val2 - ( val3 + b ) )

    w=(r*w_temp/k)
    b=(r*b_temp/k)
```

In [96]:

```
## print w and b
print('W of the implemented model is: ',w)
print('b intercept of the implemented model is: ',b)
```

```
W of the implemented model is: [-1.40249144  0.82610736 -0.75205572 -0.09542387 -1.76483671
 2.76173337
 0.17792795 -2.47815113  3.29377206 -2.52745313 -1.96668914  1.13982008
-3.3257751 ]
b intercept of the implemented model is: 22.38434562063488
```

In [97]:

```
##### Plot prediction on on Train-Data #####
y__actual = x_train['cost']
```

```


y_pred = x_train.drop('cost',axis=1).dot(w.T) + b

plt.scatter(y_actual,y_pred)
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title('Actual vs Predicted price with manual SGD on train data')
plt.show()

train_loss_manual = mean_squared_error(y_actual, y_pred)

print("Mean Squared Error: ",train_loss_manual)

```



Mean Squared Error: 19.7516997059925

In []:

In [98]:

```


##### Plot on Test-Data #####
y_actual = x_test['cost']
y_pred = x_test.drop('cost',axis=1).dot(w.T) + b

plt.scatter(y_actual,y_pred)
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title('Actual vs Predicted price with manual SGD on test data')

plt.show()

test_loss_manual = mean_squared_error(y_actual, y_pred)
print("Mean Squared Error: ",test_loss_manual)

```



Mean Squared Error: 31.520946914475672

```
In [ ]:
```

```
In [100]:
```

```
Y_train = np.array(y_train)
Y_test = np.array(y_test)
X_train = x_train.drop('cost',axis=1)
X_test= x_test.drop('cost',axis=1)

sgd = SGDRegressor(learning_rate='constant', eta0=0.01, n_iter_no_change=10000)
sgd.fit(X_train, Y_train)
```

```
Out[100]:
```

```
SGDRegressor(alpha=0.0001, average=False, early_stopping=False, epsilon=0.1,
             eta0=0.01, fit_intercept=True, l1_ratio=0.15,
             learning_rate='constant', loss='squared_loss', max_iter=1000,
             n_iter_no_change=10000, penalty='l2', power_t=0.25,
             random_state=None, shuffle=True, tol=0.001,
             validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [ ]:
```


```
In [101]:
```

```
y_actual = Y_train
y_pred = sgd.predict(X_train)

plt.scatter(y_actual,y_pred)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual vs Predicted prices with sklearn regressor on TRAIN Data')
plt.show()

train_loss_sklearn = mean_squared_error(y_actual, y_pred)

print("Mean Squared Error: ",train_loss_sklearn)
```



Mean Squared Error: 20.981205122194478

```
In [ ]:
```

```
In [102]:
```

```
y_actual = Y_test
```

```

y_pred = clf.predict(x_test)


plt.scatter(y_actual,y_pred)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual vs Predicted prices with sklearn regressor on TEST Data')
plt.show()

test_loss_sklearn = mean_squared_error(y_actual, y_pred)

print("Mean Squared Error: ",test_loss_sklearn)

```

Actual vs Predicted prices with sklearn regressor on TEST Data



Mean Squared Error: 25.604723876550317

In []:

In [103]:

```

sk_intercept = clf.intercept_[0]
print("sklearn Intercept:  \n\n",sk_intercept)

```

sklearn Intercept:

22.531274217756902

In [104]:

```

sk_weight = clf.coef_
print("sklearn weight:  \n\n", sk_weight)

```

sklearn weight:

```

[-0.84814632  0.98566697 -0.05538153  0.70848266 -1.93283664  2.68144183
 -0.00689792 -3.08696866  2.10369167 -1.39201532 -2.00587491  0.83195298
 -3.68118666]

```

In []:

In [106]:

```

from prettytable import PrettyTable
table= PrettyTable()

table.field_names = ["parameter", " manually implemented SGD regressor ", " sklearn's SGD regressor"]

table.add_row(['training MSE error',train_loss_manual,train_loss_sklearn])
table.add_row(['test MSE error',test_loss_manual,test_loss_sklearn])

```

```
table.add_row(['test MSE error', test_loss_manual, test_loss_sklearn])
table.add_row(['b intercept', b, sgd.intercept_])
print(table)
```

parameter	manually implemented SGD regressor	sklearn's SGD regressor
training MSE error	19.7516997059925	20.981205122194478
test MSE error	31.520946914475672	25.604723876550317
b intercept	22.38434562063488	[22.7555584]

In []:

In []:

In []:

In []: